

Advanced Applications of Generative AI in Actuarial Science

Dr Simon Hatzesberger

1/2 October 2026 | Vilnius, Lithuania

Lecturer



Dr Simon Hatzesberger

Actuary DAV
Certified Actuarial Data Scientist
Certified Enterprise Risk Actuary

Education



MSc in Financial
Mathematics and
Actuarial Sciences



MSc Computer Science
PhD Mathematical
Stochastics

Actuarial Associations



International Actuarial Association
Association Actuarielle Internationale

DAV & AAE: Actuarial Data Science committees
IAA: Vice-Chair, AIDE Committee

Career

Deloitte.

Actuarial & Insurance
Services, Deloitte

Allianz 

Allianz Private
Krankenversicherung
(Allianz Health Germany)

Other



EIOPA: Expert Group on Data Use in Insurance
Universities of Applied Sciences: Lecturer



Introduction & Welcome

A Quick Introduction Round (≈ 1 min/person)



Who are you?

Your name, your company and your area of work.



Your experience with AI

Pick a level from 0 to 3 on the right and add one short example if relevant.



What you expect

One thing you want to learn, solve or take back to your own work.

Your experience with AI, Generative AI or Machine Learning

0

No experience

Not used AI tools in your work so far.

1

Occasional use

Occasional use of chat assistants for single tasks.

2

Regular use

Regular chat use and first steps with an API or notebooks.

3

Advanced

Built API workflows, evaluation or agents that others rely on.

Goals and Purpose of the Seminar

Develop the practical skills and judgement to apply Generative AI effectively and responsibly in actuarial work.



Understand capabilities and limitations

Recognise where Large Language Models add value, how they fail, and which risks matter in insurance.



Choose the right approach

Assess when Generative AI is appropriate and when statistics, Machine Learning or conventional programming fit better.



Build practical workflows

Combine Large Language Models with documents and tools, and extend them into agents — each technique through a guided actuarial use case.



Evaluate results and retain control

Check output quality, trace results to sources, and build reproducibility and human review into every workflow.

Our approach Two days of conceptual foundations, five applied case studies and guided Python exercises.

Agenda

Thursday, 1 October 2026

08:45–09:00	Registration
09:00–09:15	Introduction & Welcome (EAA)
09:15–11:00	Foundations of Generative AI in Actuarial Practice
11:00–11:15	Coffee Break
11:15–12:45	Hands-on Case Study 1: Improving Claim Cost Prediction with LLM-Extracted Features from Unstructured Data
12:45–13:45	Lunch Break
13:45–15:15	Hands-on Case Study 2: Market Comparison Using Retrieval-Augmented Generation
15:15–15:30	Coffee Break
15:30–17:00	Hands-on Case Study 3: Image-based Vehicle Damage Classification and Localisation
17:30–19:00	City Tour
19:00	Dinner

Friday, 2 October 2026

09:00–10:45	Hands-on Case Study 4: Multi-Agent Systems for Actuarial Code and Data
10:45–11:00	Coffee Break
11:00–12:30	Hands-on Case Study 5: Report Generation and Quality Assurance with Generative AI
12:30–13:30	Lunch Break
13:30–15:00	Outlook and Discussion: Further Applications, Challenges, and Future Developments
15:00	Concluding Remarks, Closing of Seminar (EAA)

■ Hands-on case study

Supplementary Material

Article

Home > European Actuarial Journal > Article

Advanced applications of generative AI in actuarial science: case studies beyond ChatGPT

Case Study | Published: 10 August 2026
Volume 16, pages 481–523 (2026) [Cite this article](#)



European Actuarial Journal

Full case studies, results and risk discussion.

link.springer.com/article/10.1007/s13385-026-00464-9

Notebooks

Advanced Applications of Generative AI in Actuarial Science

Case Study: GenAI-Driven Market Comparison

[Open in Colab](#) [Open in Kaggle](#)

Description

This notebook demonstrates how Generative AI (GenAI) can streamline market comparisons in the insurance industry by extracting and harmonizing key information from unstructured annual reports. We showcase how a Retrieval-Augmented Generation (RAG) pipeline — combined with Structured Outputs — can efficiently process diverse data formats (e.g., solvency capital ratios, discount rates per duration, insurer financial strength ratings) to support data-driven decision-making. To assess robustness, the pipeline benchmarks five large language models — Claude Sonnet 4.6, GPT-4.1, GPT-4.1 mini, GPT-3.4, and GPT-3.4 mini — using deterministic scoring against ground-truth reference values with repeated runs per configuration. The approach is adaptable to various document types (e.g., risk reports, sustainability reports, insurance product comparisons) and has significant potential for automating manual actuarial tasks.

One notebook per case study (1–4).

github.com/IAA-AITF/Actuarial-AI-Case-Studies

Getting Started with the Seminar Exercises

We provide the material. You run it, inspect it and adapt it with guidance.

What is a Jupyter notebook?

An interactive document that holds explanations, Python code and the results it produces.

- 1 Read** a short explanation of the task and/or code
- 2 Run** a small block of prepared Python code
- 3 Inspect** the result, which appears directly below



Open Google Colab in your browser

Everything runs in the browser on one fixed Python setup, identical for everyone. No local installation is needed.



Start from the prepared notebooks

Sample data and worked examples give you a starting point. Core exercises run together in the room, extra tasks if time allows.



Read without a key, run with one

Every notebook ships with its outputs, readable without a key. With a key, cells call the model live; bulk runs come from labelled recordings, repeated via the **RUN_...** switches in the setup cell.


Repository and setup check

github.com/simonhatzesberger/advanced-genai-actuarial

Open `00_getting_started.ipynb` once before the first exercise.

The Shared API Key: Access and Responsible Use

An API lets the notebook send a request to a Large Language Model.
The LLM API key authorises that access and links the usage to one seminar account.



Keep the key private

Store it in Colab Secrets or a local .env file — never in a notebook cell, a screenshot or the repository.

Keep usage small

Seminar examples only, no bulk jobs. Never send company, client or personal data to the model.

Temporary access

Handed out in the room, spend-capped, and revoked once the seminar ends.

Live Demonstration

00_getting_started.ipynb



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

Foundations of Generative AI in Actuarial Practice

What You Will Take Away

From how a Large Language Model produces a token to the four techniques that make it useful in actuarial work.

- **Foundations of Generative AI**

The basic working principles of Large Language Models.

- **Providers and models**

How they differ in performance, speed, capability and multimodality.

- **Access and deployment**

Reaching models through an API, and local deployment versus the cloud.

- **Settings and parameters**

What you can configure on an API call and how it changes behaviour.

- **Prompt engineering**

Zero-shot, few-shot and chain-of-thought prompting.

- **Structured Outputs**

Constrain a response to a predefined output schema.

- **Function Calling**

Let a model invoke external tools and APIs.

- **Fine-Tuning**

Train on input–output pairs to tailor a model to a task.

- **Retrieval-Augmented Generation**

Query and retrieve from your own documents instead of guessing.


- **Hands-on practice**

Notebooks 01 and 02 apply them to actuarial examples; RAG follows in Case Study 2.

From Web Chat to API

Web chat

- Free-form and conversational
- A human in the loop — manual operator
- The interface hides parameters and response metadata
- Not scriptable, not auditable
- **Best for** prototyping and one-off questions



API

- Programmatic and scriptable
- Batch the same prompt across thousands of inputs
- Full parameter control (model, temperature, ...)
- Structured response (token usage, status)
- **Best for** automation, integration and audit

```
response = client.responses.create(  
    model=MODEL,  
    input=prompt,  
)  
print(response.output_text)
```

The chat box is for prototyping; the API is for production.

How an LLM Works — One Token at a Time



↻ Append the token to the sequence, then repeat — until a stop token or the length limit

Text becomes integer IDs from a fixed vocabulary — roughly 50,000 entries in GPT-2 small. “EIOPA” splits into pieces, “ is” is one token, and you pay per token, not per word.


Each token ID becomes a learned vector — 768 dimensions in GPT-2 small. The model places semantically related tokens near each other in that geometry.

Each token reads the earlier ones across multiple heads and layers — $12 \times 12 = 144$ in GPT-2 small. Different heads track syntax, coreference and position.


The final layer gives logits; softmax turns them into probabilities. Temperature and top-p reshape that distribution — then one token is picked.

Example of Next Token Prediction

From notebook 01 — prompt: “The Solvency Capital Requirement is ...”



Left GPT-2’s ten most likely next tokens hold only 41.4% together, led by common words (“a”, “the”, “to”).



Right Top-p keeps the most likely tokens until their probabilities add up to p, and samples only from those: at p = 0.35, 7 tokens (navy).

GPT-2 small vs. Frontier Model

Prompt: “An actuary’s main job is to ...”

GPT-2 small 124M parameters, 2019

GREEDY DECODING

“... make sure that the person is well-informed about the subject matter of the work. The actuary’s main job is to make sure that the person is well-informed about the subject matter ...”

SAMPLED, $T = 1.0$, $TOP-P = 0.9$

“... record and analyze a deceased person’s life. If someone is deceased, the funeral director would typically file it under a specific actuary’s name ...”

Fluent, and wrong: repetition loops and confident invention.

Frontier model e.g. GPT-6 Sol, 2026

REPRESENTATIVE OUTPUT

“... price risk, set reserves for future claims, and report the resulting liabilities under the relevant regulatory framework (Solvency II in the EU, IFRS 17 globally). The day-to-day work blends statistical modelling, business judgment, and explanation to non-technical stakeholders.”

Same mechanism, vastly more scale and post-training.

Capabilities of LLMs

Frontier models

Start with gpt-5.4 for complex reasoning and coding, or choose gpt-5.4-mini and gpt-5.4-nano for lower-latency, lower-cost workloads.

gpt-5.4

GPT-5.4 New
Best intelligence at scale for agentic, coding, and professional workflows

Model ID: gpt-5.4

Reasoning: none low medium high xhigh

Input price: \$2.50 / Input MTok

Output price: \$15 / Output MTok

Latency: Fast

gpt-5.4-mini

GPT-5.4-mini
Our smallest and most efficient model

Model ID: gpt-5.4-mini

Reasoning: none low medium high xhigh

Input price: \$0.50 / Input MTok

Output price: \$3 / Output MTok

Latency: Fast

gpt-5.4-nano

GPT-5.4-nano
Our smallest and most efficient model

Model ID: gpt-5.4-nano

Reasoning: none low medium high xhigh

Input price: \$0.10 / Input MTok

Output price: \$0.50 / Output MTok

Latency: Fast

Specialized models

Purpose-built for specific tasks.

Image Models for image generation and editing	GPT Image 1.5 State-of-the-art image generation model	gpt-image-1-mini A cost-efficient version of GPT Image 1
Realtime Models for realtime speech-to-speech	gpt-realtime-1.5 The best voice model for audio in, audio out	gpt-realtime-mini A cost-efficient version of GPT Realtime
Speech generation Models for generating natural-sounding speech from text	GPT-4o mini TTS Text-to-speech model powered by GPT-4o mini	
Transcription Models for transcribing speech into text	GPT-4o Transcribe Speech-to-text model powered by GPT-4o	GPT-4o mini Transcribe Speech-to-text model powered by GPT-4o mini

Features

- Streaming** Supported
- Structured outputs** Supported
- Distillation** Supported
- Function calling** Supported
- Fine-tuning** Not supported

Tools

Tools supported by this model when using the Responses API.

- Web search**
- File search** Supported
- Code interpreter** Supported
- Apply patch** Supported
- Computer use** Supported
- Tool search** Supported

Comparison of Selected LLMs

Provider	Model	Tier	\$ / 1M tok, in / out	Context	Reasoning
OpenAI	GPT-6 Astra	Frontier	10.00 / 50.00	1M	yes
OpenAI	• GPT-6 Sol	Balanced	2.00 / 10.00	1M	yes
OpenAI	• GPT-6 Luna	Cheapest	0.10 / 0.50	1M	yes
Anthropic	Claude Fable 5.1	Frontier	10.00 / 50.00	1M	yes
Anthropic	Claude Sonnet 5	Balanced	2.00 / 10.00	1M	yes
Anthropic	Claude Haiku 4.5	Cheap + fast	1.00 / 5.00	200K	yes
Google	Gemini 3.1 Pro	Frontier	2.00 / 12.00	1M	yes
Google	Gemini 3.5 Flash-Lite	Cheap + fast	0.30 / 2.50	1M	yes
Mistral	Mistral Small 4	Open weights	0.15 / 0.60	256K	yes
Meta	Muse Spark 1.3	Frontier	1.25 / 4.25	1M	yes
Meta	Llama 4 Scout	Open weights	self-hosted	10M	no

Standard list prices in USD per 1M tokens, from each provider's own pricing page, checked 22 September 2026; Gemini 3.1 Pro doubles above a 200K-token prompt.

• = called by this seminar's notebooks, which also use GPT-5.4, GPT-4o (and a fine-tune), text-embedding-3, gpt-image-2.5 and, locally, GPT-2 and Qwen2.5-0.5B.

Cloud vs. Local

Dimension	Cloud API (OpenAI / Anthropic / ...)	Local, self-hosted (Ollama, vLLM, ...)
Latency	100 ms to a few seconds from the EU	Depends on hardware; GPU fast, CPU slow
Cost shape	Per token; no fixed cost; predictable	Hardware, ops and electricity; predictable at scale
Capability ceiling	Frontier (200B+ dense, MoE, reasoning)	Open weights lag the frontier by 6–12 months on hard tasks
Data residency / GDPR	Goes to the provider's servers; needs a DPA or SCCs	Stays on your hardware, period
Confidentiality / IP	The provider sees prompts; check the terms on training	Only you see the prompts
Operations burden	None — the provider scales, patches and updates	High — you scale, patch and monitor GPUs
Model availability	The provider's catalogue; deprecations happen	You pin a checkpoint forever
Best fit	Most workflows; rapid iteration; sensible vendor risk	Confidential or regulated data; on-prem mandates; very high volume

The GPT-2 demo in notebook 01 is a teaching toy, not a production local stack.

For real local deployment: **Ollama with Qwen 3.8-27B or Mistral Small 4 on a GPU.**

Selected API Parameters at a Glance

Parameter	What it does	When to touch it
<code>model</code>	Which model version to call	Pick the cheap tier first; upgrade only if it fails on real examples
<code>instructions</code>	System prompt — persistent role, audience, format and tone	Almost always: the single biggest lever on output shape
<code>max_output_tokens</code>	Length cap, holding down cost and latency	Whenever you know the response should be short
<code>temperature, top_p</code>	Sampling knobs, on models that expose them	See them on GPT-2 small in notebook 01; recent API models accept them only without reasoning, so the case studies set at most temperature 0

Selected Prompt Engineering Techniques

Zero-shot

Ask the question. No examples.

Claim note: "Rear-ended at a red light; the other driver admitted fault. Bumper and tailgate dented."

Give the cause of loss as one category.

USE WHEN

The task is common and the model's defaults will do.

Best for common tasks and quick drafts.

Few-shot

Give the model a few worked examples first.

Classify the cause of loss of a motor claim note.

"Hail dented the roof." -> weather
"Car gone overnight." -> theft
"Rear-ended at a red light; other driver at fault." ->

USE WHEN

Zero-shot comes back fuzzy on format or style.

Best for shaping format, classification, constraining style.

Chain-of-thought

Ask the model to reason step by step.

Claim note: "Rear-ended at a red light; the other driver admitted fault. Bumper and tailgate dented."

Reason step by step about the cause of loss, then end with "Cause: <category>".

USE WHEN

Multi-clause logic or multi-stage extraction. (Kojima et al., NeurIPS 2022)

Caveat: plausible-sounding reasoning is not the same as correct.

Limitations of LLMs

Hallucination

Plausible-sounding text that is factually wrong — common on rare facts, specific numbers, names and citations.

Prompt injection and jailbreaks

User content can override your instructions. Always quote, delimit, and treat data as data.

Training cut-off staleness

The model knows events only up to its training data. Live data — today's policy database, yesterday's market — is invisible.

Cloud confidentiality

Cloud APIs see your prompts. Check provider terms on training, retention and region; prefer local for sensitive personal data.

Prompt sensitivity

Small wording changes can change the output. Determinism is best-effort, never guaranteed.

Evaluation is hard

There is no RMSE for an LLM output. You need spot checks, golden sets or judge models.

Live Demonstration

01_llm_basics.ipynb



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

Motivation for Advanced Techniques



Chatbot limitation

GenAI is used only in chat interfaces, never inside a process.

LLMs through APIs



Free-text chaos

Responses vary from run to run and are hard to process further.

Structured Outputs



No real actions

The model can suggest, but it cannot execute anything itself.

Function Calling



Low accuracy

Standard models are not precise enough for specialised tasks.

Fine-Tuning



Missing context

It knows nothing of your policies, data, documents or recent events.

Retrieval-Augmented Generation

The rest of this block One slide per technique; notebook 02 runs three of them hands-on, Case Study 2 the fourth.

Structured Outputs

Idea Constrain the response to a (Pydantic-typed) JSON schema — no fragile parsing.

WHAT IT DOES

Pass a Pydantic class through `text_format`. The API holds the reply to its JSON schema, and `output_parsed` arrives as a validated Python object.

WHY IT MATTERS

Same fields and types on every call, so code can compute with the answers and check them. Nested schemas hold what a flat one cannot: coverage lines, endorsements, discounts.

ACTUARIAL EXAMPLE

Ten synthetic car-insurance contracts on the Quebec form (RISCBAC), in three stages: free text, a flat typed schema, then a nested schema with an enum, booleans, optional fields and an ordered cover level. Code checks then flag every contract the model got wrong.

HOW TO CALL IT

```
contract = client.responses.parse(  
    model=MODEL,  
    input=policy_text,  
    instructions=CONTRACT_INSTRUCTIONS,  
    text_format=CarInsuranceContract,  
).output_parsed
```

Function Calling

Idea Expose Python code as tools — the model decides when to call deterministic logic.

WHAT IT DOES

Describe each function with a name and a JSON schema in `tools=`.
Four steps: the model asks for a call, your code runs it, the result goes back under the call id, the model answers.

WHY IT MATTERS

The numbers come from your database, not from the model's memory.
The tools only read, take checked arguments and run fixed SQL; every call, argument and query is logged.

ACTUARIAL EXAMPLE

The PKV-Sterbetafel 2022–2025 in an SQLite database, behind two read-only tools: the model compares 2022 with 2025 for men aged 85 to 100 and writes a table and a short trend analysis. Code checks the numbers: 16 of 16 rows match.

HOW TO CALL IT

```
response = client.responses.create(  
    model=MODEL,  
    input=messages,  
    tools=TOOLS,    # two read-only tools  
)
```

Fine-Tuning

Idea Continue training on labelled examples — permanent domain adaptation.

WHAT IT DOES

Collect prompt-and-response pairs as JSONL, upload them, launch a training job and receive a new model ID. Inference runs through the same API — the ID carries the adapted weights.

WHY IT MATTERS

Bakes in domain terminology (e.g., Solvency II, IFRS 17) and keeps decisions consistent across runs. Small fine-tuned models often match frontier accuracy on narrow tasks at a fraction of the cost.

ACTUARIAL EXAMPLE

Notebook 02: LoRA on the open-weights Qwen2.5-0.5B-Instruct, 20 claim notes, one cause-of-loss label. Case Study 3 calls a GPT-4o fine-tuned on 900 car-damage photos.

HOW TO CALL IT

```
# notebook 02: Qwen2.5-0.5B-Instruct, trained on CPU
from peft import LoraConfig, get_peft_model
cfg = LoraConfig(r=8, lora_alpha=16,
                 target_modules=["q_proj", "v_proj"])
model = get_peft_model(base_model, cfg)
# 540,672 of 494.6M parameters train (0.11%)
```

Retrieval-Augmented Generation

Idea Ground the model in your own documents — and close the knowledge-cutoff gap.

WHAT IT DOES

Five stages: index the documents (chunk, embed, store), vectorise the query, retrieve the most similar chunks, augment the prompt with them, then generate the answer — typically with structured outputs.

WHY IT MATTERS

Current, because sources update without retraining. Accurate, because answers are grounded in the cited chunks. Flexible, because document sources can be added or swapped at any time.

ACTUARIAL EXAMPLE

Extract three KPIs from three insurers' 2025 annual reports: the solvency capital ratio, the euro discount-rate curve by duration, and the financial-strength ratings by agency.

HOW TO CALL IT

```
client.embeddings.create(  
    input=chunk, model=...)   
cosine_similarity(query_vec, docs)   
client.responses.parse(...)
```

Live Demonstration

02_llm_techniques.ipynb



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

Hands-on Case Study 1


Improving Claim Cost Prediction with LLM- Extracted Features from Unstructured Data

Claim Descriptions Carry Cost Drivers

Age	Gender	Weekly wage	Hours	Claim description	Ultimate cost
36	M	296	38	LIFTING HEAVY WIRE HERNIA LEFT GROIN MUSCLE	10,534
32	M	200	70	PARTICLES WHILST GRINDING IRRITATION LEFT EYE	416
40	M	529	38	UNLOADING PALLET STRAINED LEFT TENDON LEFT HAND	20,498
26	M	672	40	REDBACK SPIDER BITE RIGHT FOOT RIGHT FRACTURE	2,530
31	F	441	40	DERMATITUS RASH RIGHT WRIST INJURY	30,096
20	M	505	38	CAUGHT IN PENETROMETER CRUSHED LEFT INDEX FINGER LEFT HAND	5,872

3,000 claims, each with 12 tabular fields and one description of at most 13 words (median 7), in capitals, with abbreviations and typos.
2,420 descriptions are distinct. Costs in USD.

What the Model Has to Predict



The target

Ultimate incurred cost per claim: median USD 3,573, mean 10,685, maximum 428,384. A quarter of the claims stay below USD 1,000.

Why the log

On the raw scale a few large claims dominate every error measure. On the log scale the distribution is compact.


The inputs

18 tabular features: e.g. age, gender, marital status, dependants, weekly wage, hours, days, part or full time, accident date and hour, reporting delay.

Baseline: predict log cost from the 18 tabular features. **Enhanced:** add what the LLM reads from the description.

Machine Learning in One Slide

Supervised learning



Gradient boosting



Many small decision trees, each fitted to the errors the previous ones left. The prediction is their sum. Strong on tabular data, little tuning needed.

Train and test



The model is fitted on the training claims only. The test claims are scored once at the end, as stand-ins for next year's claims. A model that memorises its training data looks good there and fails here.

Stratified by cost decile, so that both parts hold the same mix of small and large claims.

Cross-validation repeats the split four times inside the training data, to tune the model and to see how much the result moves. It is an exercise in the notebook.

Same data, same split, same model type for baseline and enhanced: the only difference is the extra features.

How Good Is a Prediction?

	Actual	Predicted	Log error	Squared
Claim 1	2,000	3,000	+0.41	0.16
Claim 2	10,000	6,000	-0.51	0.26
Claim 3	500	1,500	+1.10	1.21

MAE = mean of $|\log \text{ error}| = 0.67$

RMSE = $\sqrt{\text{mean of squared errors}} = 0.74$; it punishes the big miss of claim 3 harder

Miss factor = $\exp(\text{MAE}) = 1.96$: a typical prediction is off by about a factor of 2.0

Three measures, all on log cost

RMSE: root mean squared error; large misses weigh most. Lower is better.

MAE: mean absolute error; every miss counts in proportion. Lower is better.

R²: share of the variation in log cost the model explains: 0 means no better than the average, 1 means perfect.

In the article, the LLM features took the miss factor on the test claims from 2.94 to 2.30.

Numbers invented for illustration; the case study reports RMSE, MAE and R² on the 600 test claims.

Fixed Classes Instead of Free Text

Body part (8)	for example
HEAD_FACE	head, face, ear, jaw, teeth
EYES	eye, cornea, eyelid
NECK	neck, cervical spine, whiplash
TORSO	back, chest, ribs, abdomen, groin
UPPER_EXTREMITY	shoulder, arm, elbow, forearm, wrist
HAND_FINGERS	hand, palm, thumb, fingers
LOWER_EXTREMITY	hip, leg, knee, ankle, foot, toes
OTHER	skin, whole body, not stated

```
class ClaimFeatures(BaseModel):
    body_part_category: Literal[...] # 8 names
    cause_of_injury_category: Literal[...] # 13 names
    number_of_body_parts_injured: int = Field(ge=0)
```

Cause (13)	for example
LIFTING_CARRYING	lifting, carrying, pushing, pulling
IMPACT	struck by or against, falling object
FALL_SLIP_TRIP	fall, slip, trip, stumble
STRAIN_SPRAIN	bending, twisting, overreaching
LACERATION	cut by a sharp object, scraping
CHEMICAL_BURN	heat, steam, sparks, chemicals
COMPRESSION	caught in or between, crushed
CONTACT	foreign body, shock, bite, assault
ERGONOMIC	ordinary tool use, no other cause
WALKING	walking, stepping, climbing, no fall
DRIVING	driving, road accident
MEDICAL	condition named, no mechanism
OTHER	no cause stated, or none of these

The model can only answer with one of these names: no mapping step afterwards, no pattern rules to maintain.

One Claim, One Call

CAUGHT BETWEEN GRINDER HAND TOOL CUT RIGHT THUMB



```
{"body_part_category": "HAND_FINGERS",  
  "cause_of_injury_category": "COMPRESSION",  
  "number_of_body_parts_injured": 1}
```

1,314 tokens in, 37 out: USD 0.00015 per claim, 1.5 cents per hundred.
The instructions (704 words, the class definitions and rules) travel with every call.

The rules at work

- **Primary body part:** the thumb. HAND in HAND TOOL names a tool, so it neither sets the class nor counts.
- **Mechanism over outcome:** CAUGHT BETWEEN (COMPRESSION) wins over CUT (LACERATION).
- **Written down once:** the same rules for every claim, changed in one place. Other rules would be defensible; what matters is that they are explicit.

All 3,000 claims: 2,420 distinct descriptions, one call each, 16 in parallel: 196 seconds and USD 0.37.

What the Model Made of the Descriptions

Description	Body part	Cause	Parts
INJURED WRIST	UPPER_EXTREMITY	OTHER	1
LOWER BACK PAIN LOWER BACK STRAIN	TORSO	STRAIN_SPRAIN	1
MOTOR VEHICLE ACCIDENT MOTOR VEHICLE ACCI	OTHER	DRIVING	0
CARING A STRAIN RIGHT FOREARM	UPPER_EXTREMITY	OTHER	1
LIFTING QUAD FROM UTE LACERATION RIGHT HAND	HAND_FINGERS	LIFTING_CARRYING	1
REDBACK SPIDER BITE RIGHT FOOT RIGHT FRACTURE	LOWER_EXTREMITY	CONTACT	1
SPRAINED RIGHT ANKLE FRACTURE RIGHT ELBOW	LOWER_EXTREMITY	STRAIN_SPRAIN	2
SLIPPED ON ROLLER TENDONITIS RIGHT SHOULDER	UPPER_EXTREMITY	FALL_SLIP_TRIP	1

Checks without labels

Asked twice: 99 of 100 answers identical in all three fields.


Live vs recorded: 487 of 500 identical.

More reasoning: effort “low” changed 15 of 100 answers and cost 1.8 times as much.

Read closely: CARING is probably carrying, yet the model chose OTHER.

Without reference labels you can measure consistency; whether the classes are useful, the cost model shows.

What the Classes Say About Cost



Torso claims cost 15 times as much as eye claims at the median: USD 7,038 against 468.

Causes split the same way: lifting, strains and road accidents above 6,500; impacts, cuts, burns and contact below 1,200.

The count shows no clear pattern.

The classes separate cheap from expensive claims; whether they add to the tabular features is the next question.

Results: Baseline Versus Enhanced

Test set, 600 claims	Baseline	Enhanced	Improvement	Article (GPT-4o mini + mapping)
RMSE	1.3425	1.1216	16.5%	1.345 → 1.101
MAE	1.0729	0.8452	21.2%	1.078 → 0.834
R ²	0.2274	0.4607	+0.233	0.224 → 0.481
Miss factor exp(MAE)	2.92	2.33		2.94 → 2.30

Baseline: 18 tabular features. **Enhanced:** the same plus GPT-6 Luna's three fields, one-hot encoded (38 features). Same split, same model type, settings from the article.

A typical prediction is off by a factor of 2.3 instead of 2.9. Close to the article's gain (16.5% against about 18%) with a shorter pipeline: no mapping step, no pattern rules. Not better: the article's route reached 1.101.

Cross-validation in the notebook's exercise confirms it: 1.323 ± 0.020 against 1.123 ± 0.017 , corrected t-test $p = 0.0005$.

Three fields read from one line of text cut the error by a sixth and double the explained variance.

Which Fields Carry the Gain

Features	Count	Test RMSE	vs baseline
Baseline	18	1.3425	–
+ number of body parts	19	1.3415	–0.001
+ cause	30	1.1892	–0.153
+ body part	25	1.1452	–0.197
+ all three	38	1.1216	–0.221

Body part does most of the work, cause adds to it, the count adds nothing. As in the article. One test split: read the order, not the third decimal.

Feature importance, enhanced model

Top five	Share
Weekly wage	0.232
Body part: TORSO	0.132
Reported year	0.085
Body part: HAND_FINGERS	0.067
Cause: IMPACT	0.061

By group: tabular 0.56, body part 0.25, cause 0.17, count 0.02. The extracted fields carry 44% of the model's importance.

The model uses what it is given: where the body part is known, it leans on it heavily.

Small Claims, Large Claims

Ultimate cost, USD	Claims	RMSE baseline	RMSE enhanced
below 1,000	155	1.606	1.095
1,000 to 5,000	194	0.740	0.801
5,000 to 20,000	179	0.985	0.774
above 20,000	72	2.362	2.164

Better in three bands of four, most at the small end. Between 1,000 and 5,000 the enhanced model is slightly worse, as in the article.

Large claims stay hard: above 20,000 the enhanced model still under-predicts by a factor of about 6 on average. One line of text does not reveal a long recovery.

Check the error where the money is: an average improvement can hide the band that matters for reserving.

Exercises: Your Own Feature, a Fair Comparison

(a) Design your own feature

Worked example: nature of injury (fracture, sprain, cut, bruise, burn ...), a fourth field in the schema, extracted for all claims for USD 0.16.

Test RMSE 1.122 → 1.044. Cross-validated the gain is smaller and not significant for RMSE ($p = 0.07$), significant for MAE ($p = 0.04$).

Your turn: severity, third party involved, lost time ...

(b) Cross-validation, done properly

Four folds, twice: once to choose trees and depth, once with other folds to measure.

CV RMSE 1.323 ± 0.020 against 1.123 ± 0.017 .

Fold scores share training data, so the plain t-test is too optimistic. The corrected test (Nadeau and Bengio) still gives $p = 0.0005$.

A new feature has to pay its way on data the model has not seen, with an honest test.

Live Demonstration

`03_cs1_claims_text_features.ipynb`



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

What This Means for Actuarial Work

Text is a data source

Claim descriptions, adjuster notes and emails hold cost drivers that tabular data misses. Three fields from one line doubled the explained variance.

Fix the classes first

Deciding which classes exist and which rule wins is actuarial work. A schema makes the model stick to it.

Measure what you can

Without reference labels, check consistency and let the downstream model show the value. With labels, score the extraction too.

Keep the pipeline auditable

Fixed model version, saved extractions, written rules. Re-extract and compare before a model update reaches production.

The same pattern works for underwriting notes, fraud indicators and complaint texts.

Hands-on Case Study 2

Market Comparison Using Retrieval-Augmented Generation

Comparing Insurers Means Reading Hundreds of Pages

One figure, the 2025 solvency ratio, as each group reports it:

AXA


Universal Registration Document
2025
588 pages · PDF p. 45

As of December 31, 2025, the Group's Eligible Own Funds ("EOF") amounted to €56.4 billion and the Solvency II ratio to 224%, compared to €55.9 billion and 216%, respectively, as of December 31, 2024.

In running text, next to own funds and last year's figures.

Generali

Annual Integrated Report 2025
458 pages · PDF p. 12



A dashboard tile with the change against 2024. A footnote calls the ratio preliminary.

Zurich

Annual Report 2025
430 pages · PDF p. 6



Swiss Solvency Test instead of Solvency II, estimated as of 31 December 2025.

Three reports, 1,476 pages, no common layout. Collecting such figures by hand is slow and error-prone.

Three Targets, Three Kinds of Data

ONE NUMBER, ONE LABEL

Solvency ratio

Ratio in percent plus the regime:

AXA 224% and Generali 219% (Solvency II), Zurich 259% (SST).

A TABLE OF VARYING LENGTH

EUR discount curve, 31 Dec 2025

Discount rate by duration for insurance liabilities.

AXA reports 10 durations, Generali 18, Zurich 5.

A LIST OF RECORDS

Financial strength ratings

Agency, rating and outlook, three per group.

Zurich: S&P AA, Moody's Aa2, AM Best A+ Superior, all stable.


IFRS 17 discounting curve ZC - Contract portfolios measured under non - VFA

Currency	EUR		CHF		CZK		CNY	
Maturity (years)	1/12/2025	31/12/2024	31/12/2025	31/12/2024	31/12/2025	31/12/2024	31/12/2025	31/12/2024
1	2.22%	2.47%	-0.08%	0.02%	3.56%	3.83%	1.84%	1.56%
2	2.30%	2.32%	0.00%	-0.01%	3.63%	3.70%	1.86%	1.64%
3	2.42%	2.32%	0.09%	0.03%	3.73%	3.69%	1.88%	1.69%
4	2.53%	2.35%	0.19%	0.09%	3.82%	3.70%	1.99%	1.80%
5	2.62%	2.37%	0.28%	0.14%	3.90%	3.72%	2.14%	1.92%
6	2.70%	2.40%	0.36%	0.18%	3.97%	3.75%	2.20%	2.01%
7	2.79%	2.43%	0.44%	0.22%	4.04%	3.79%	2.25%	2.10%
8	2.86%	2.45%	0.50%	0.26%	4.10%	3.82%	2.30%	2.16%
9	2.93%	2.47%	0.57%	0.30%	4.16%	3.85%	2.33%	2.18%
10	3.00%	2.50%	0.63%	0.35%	4.22%	3.89%	2.37%	2.20%
15	3.25%	2.56%	0.91%	0.63%	4.40%	3.99%	2.70%	2.42%
20	3.35%	2.49%	1.13%	0.88%	4.39%	3.98%	2.82%	2.53%
25	3.39%	2.51%	1.30%	1.08%	4.30%	3.92%	2.84%	2.50%
30	3.41%	2.58%	1.43%	1.24%	4.19%	3.86%	2.93%	2.56%
35	3.41%	2.65%	1.54%	1.37%	4.10%	3.80%	3.13%	2.79%
40	3.40%	2.71%	1.62%	1.48%	4.01%	3.75%	3.45%	3.13%
45	3.40%	2.77%	1.69%	1.56%	3.94%	3.71%	3.75%	3.47%
50	3.39%	2.82%	1.75%	1.63%	3.88%	3.67%	4.00%	3.75%

Generali, PDF p. 240. Four currencies and two dates in one table; the model has to pick the EUR column for 31/12/2025.

45 reference values, copied by hand from the reports: 3 solvency ratios, 33 curve points, 9 ratings.

The 3-Stage Approach



1

Preprocessing once per report

Read the PDF page by page, clean the text and cut it into chunks of 2,000 characters that overlap by 300. `text-embedding-3-large` turns each chunk into a vector; the vectors stay in memory.

2

Prompt augmenting once per question

Embed the question the same way, rank all chunks by cosine similarity and paste the ten best into the prompt, provided their similarity is at least 0.30.

3

Response generation once per question and model

The LLM answers from those chunks and returns JSON that has to match a Pydantic schema.

Stages 1 and 2 are identical for every model; only the LLM in Stage 3 changes. That keeps a model comparison fair.

Stage 1: From 1,476 Pages to 2,822 Chunks

Report	Pages	Characters	Chunks
AXA	588	1,969,238	1,139
Generali	458	1,515,508	885
Zurich	430	1,377,261	798
Total	1,476	4,862,007	2,822

Chunks: 2,000 characters each, every chunk starting 1,700 characters after the previous one, so neighbours share 300.

Vectors: $2,822 \times 3,072$ numbers from `text-embedding-3-large`. 1.34 million tokens, USD 0.17, 15 seconds, once. Later runs read them from disk.

Lost on the way: 21 pages without a text layer. `PyMuPDF` reads text, not pixels; a scanned table would need OCR.

AXA-0024, PDF p. 21: the ratings table after extraction

```

S&P Global Ratings
October 3, 2025
A+
AA-
Positive
A+
A-1+
Moody's Investors Service
October 8, 2025
Aa2
Aa2
Stable
Aa3
P-1
...

```

One cell per line. The grid with its two-level header is gone, and with it the information which column a rating belongs to.

The whole AXA report is about 470,000 tokens. The prompt that answers one question from it: about 5,200.

Stage 2: Ranking Chunks by Similarity

What a cosine similarity means

Solvency passages from the three groups, in different words: **0.60 to 0.67** to each other.

A biodiversity passage from the same reports: **0.46 to 0.51** to them.

Relevance is a modest lift over a high baseline, so retrieval ranks the chunks instead of trusting an absolute cut-off.

The threshold barely filters

Of AXA's 1,139 chunks, 1,033 clear 0.30 for the solvency question and 25 for the ratings question. The top-10 rule does the selecting.

The Top-10 chunks retrieved for AXA's ratings

Rank	Similarity	PDF page	Content
1	0.466	21	the ratings table itself
2	0.396	21–22	rest of the ratings page
3	0.364	20–21	the table's column headings
4	0.358	304–305	risk factor: liquidity and ratings
5	0.356	341	reinsurers without a rating
6	0.344	305	ratings can be withdrawn
7	0.341	340–341	collateral rules for derivatives
8	0.323	285–286	biodiversity
9	0.320	316–317	risk factors: capital
10	0.317	151–152	sustainability indicators

One prompt, two jobs: it is embedded to find the chunks and then sent to the model with them.

Stage 3: One Call, One Validated Object

The article's schema and the call

```
class FinancialStrengthRating(BaseModel):
    rater: str
    rating: str
    outlook: Optional[str]

class FinancialStrengthRatingsResult(BaseModel):
    ratings: List[FinancialStrengthRating]

response = client.responses.parse(
    model="gpt-6-luna",
    instructions=SYSTEM_PROMPT,
    input=augmented_prompt,
    text_format=FinancialStrengthRatingsResult)
```

What came back for AXA

```
{
  "ratings": [
    {
      "rater": "S&P Global Ratings",
      "rating": "AA-",
      "outlook": "Positive"
    },
    {
      "rater": "Moody's Investors Service",
      "rating": "Aa2",
      "outlook": "Stable"
    },
    {
      "rater": "A.M. Best Rating Services",
      "rating": "aa Superior",
      "outlook": "Stable"
    }
  ]
}
```

Valid JSON, and wrong: “aa Superior” is AM Best’s rating of AXA’s senior debt. 5,310 tokens in, 1,566 out, 14 seconds, USD 0.0013.

The same prompt again, with the eight others in parallel: 12 seconds, under a cent, all 45 values right.

A schema fixes the shape of the answer, not its content. Whether a value is right has to be checked.

Results: Numbers Are Easy, Ratings Are Not

Exact match, no partial credit

Solvency: ratio and regime both right.

Discount curve: every duration present, each rate right to two decimals.

Ratings: all three agencies right after agency names and spellings are normalised.

GPT-6 Luna in the walkthrough: 45 of 45 values, 9 of 9 targets.

Pass rate in %, per target

Model	Runs	Solvency	Discount	Ratings
GPT-6 Luna	3	100	100	77.8
GPT-5.4 mini	3	100	100	66.7
GPT-5.4	3	100	100	66.7
Article: GPT-5.4 mini	20	100	100	71.7
Article: GPT-5.4	20	100	100	100
Article: GPT-4.1 mini	20	100	100	66.7
Article: GPT-4.1, Sonnet 4.6	20	100	100	100

27 extractions per model cost about USD 0.004 on GPT-6 Luna, 0.06 on GPT-5.4 mini and 0.21 on GPT-5.4 on a first run. With three runs, one extraction moves a rate by 11 points.

Every failure was the same table: AXA's ratings. Generali and Zurich passed in every run, for every model.

Where It Goes Wrong: Financial Strength Ratings

1.5.1 Insurer financial strength and counterparty credit ratings

At the date of this Document, the relevant ratings for the Company and its principal insurance subsidiaries were as follows:

Agency	Date of last review	Insurer financial strength ratings			AXA's credit ratings	
		AXA SA	AXA's principal insurance subsidiaries	Outlook	Senior debt of the Company	Short-term debt of the Company
S&P Global Ratings	October 3, 2025	A+	AA-	Positive	A+	A-1+
Moody's Investors Service	October 8, 2025	Aa2	Aa2	Stable	Aa3	P-1
A.M. Best Rating Services	October 9, 2025	A+ Superior	-	Stable	aa Superior	-

AXA, PDF p. 21. Blue: holding company and insurance subsidiaries, which differ for S&P (A+ vs AA-). Dashed: debt ratings, which the question does not ask for.

The reference is a judgement, too: it takes the subsidiaries' rating where there is one and AXA SA's where there is not. A model can only follow that rule if the prompt states it.

What happened

The table was the top-ranked chunk every time. The models then read the wrong column: GPT-5.4 mini took S&P's A+ for AXA SA in all three runs; GPT-5.4 in all three, and GPT-6 Luna in two, took AM Best's "aa Superior", a senior-debt rating.

The fix we tested

The research notebook's tighter schema, with fixed agency names and a description for every field, took GPT-5.4 from 0 of 3 to 3 of 3. GPT-5.4 mini stayed at 0 of 3; GPT-6 Luna fell from 1 of 3 to 0 of 3.

Generation error with its root in Stage 1: flattening the PDF threw away the column structure the model needed.

Live Demonstration

`04_cs2_rag_market_comparison.ipynb`



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

The Same Pipeline as an App

Live-Demo

Step 5 — Generate Structured Answer

Stage 3 — Response generation. Only this stage depends on the model you pick; the retrieval above stays fixed. Each company's augmented prompt goes to the model with Structured Outputs, so the reply is JSON that validates against the schema — no regular expressions or string parsing downstream.

Generation model

- ☒ GPT-6 Luna (gpt-6-luna) — default · rejects temperature, so none is sent
- ☐ GPT-5.4 mini (gpt-5.4-mini-2026-03-17) — temperature 0, as in the article
- ☐ GPT-5.4 (gpt-5.4-2026-03-05) — temperature 0, as in the article

Generate with GPT-6 Luna

Results — GPT-6 Luna (gpt-6-luna)

AXA

224% under Solvency II

4,942 tokens in, 51 out · 1.6 s · no temperature sent

> Show raw JSON

Generali

219% under Solvency II

4,956 tokens in, 52 out · 2.1 s · no temperature sent

> Show raw JSON

Zurich

259% under SST

5,000 tokens in, 48 out · 1.5 s · no temperature sent

> Show raw JSON

Six steps on one page


- 0 Reports and key
- 1 Chunks and their vectors
- 2 Prompt and Pydantic schema
- 3 Retrieved chunks with similarity bars; top-N and threshold on sliders
- 4 The augmented prompt and its token count
- 5 One structured answer per company. Switching the model reruns only this step.

```
cd apps/cs2_market_comparison
streamlit run app.py
```

Step 5 of the app: one Structured Outputs call per report, all three in parallel.

Same prompts, schemas and retrieval settings as the notebook. One run for all three groups takes 15k input tokens.

Extra Exercise: Embeddings as Rating Factors



The paper's claim

With 500 training policies, a GLM on engineered features scores a test Poisson deviance of 1.208, the best embedding GLM 0.3591. With 500,000 policies the best embedding GLM still leads, 0.3120 against 0.3211.

The language model only supplies covariates. The GLM prices.

Our rerun: test mean Poisson deviance, 20,000 policies

Model	n = 500	n = 2,000	n = 10,000
Flat rate	0.3279	0.3282	0.3285
GLM, raw features	0.9586	0.5056	0.3234
Ridge GLM, raw features	0.3279	0.3268	0.3222
Embedding GLM, K = 12	0.3508	0.3305	0.3267
Embedding GLM, K = 48	no fit	0.3756	0.3272

Against a plain GLM the paper's pattern holds, but a flat rate beats every embedding GLM at 500 & 2,000 policies: the small-sample gain is mostly regularisation. At 10,000 the embeddings beat the flat rate.

What This Means for Actuarial Work

Retrieval sets the ceiling

Chunking, embedding model and prompt wording matter as much as the choice of LLM. Better retrieval often pays off more than a bigger model.

Say which figure you mean

When a report holds several similar figures, such as group and subsidiary ratings, the prompt or the schema has to name the one you want. That is a business decision.

Plan for the missing answer

Let the model say when the context does not contain the value, and handle that case in code instead of accepting a guess.

Keep a reviewer in the loop


Show the retrieved passages and page numbers next to each value, and compare with last year's figures before anything goes into a report.

The same pipeline works for risk reports, sustainability reports and tariff comparisons.

Hands-on Case Study 3

Image-based Vehicle Damage Classification and Localisation

From Claim Photo to Claim Record



The case study asks: Can a vision LLM, fine-tuned on 900 labelled photos, classify car damage as well as a purpose-built image model, and add the context on top?

Six Damage Classes, 7,200 Photos

Class	Photos	Share	Test set
scratch	2,349	32.6%	104
dent	2,079	28.9%	80
glass shatter	1,185	16.5%	50
lamp broken	882	12.3%	37
tire flat	534	7.4%	22
crack	171	2.4%	7
Total	7,200	100%	300

Sample: 1,500 photos, split 900 for training, 300 for validation and 300 for testing, stratified by class.

The catch: seven test photos of cracks. One photo more or less moves the crack recall by 14 points.



crack



scratch



tire flat



dent



glass shatter



lamp broken

Small, noisy, often cropped photos, some in black and white: harder than they look.

One label per photo, even where a photo shows a dent and a broken lamp at once.

Three Ways to Classify a Photo

	Image classifier (CNN) EfficientNet-B0, pretrained on ImageNet	Vision LLM, zero-shot GPT-4o with a prompt	Vision LLM, fine-tuned GPT-4o trained on the 900 photos
How it learns	Retrained on the 900 labelled photos	Not at all: the prompt lists the six classes	900 labelled photos on the provider's platform
What comes out	One of six labels	A label that fits a schema; location or severity on request	The same, tuned to these six classes
What it needs	Deep-learning know-how; a GPU helps	A careful prompt	Labelled examples and a platform that still offers fine-tuning
Where it runs	Anywhere, including on premises	The provider's API	The provider's API, for the owner's organisation only

The image classifier returns a label. The vision LLMs can say more, and every extra field is one more thing to check.

One Photo, One Call

The photo travels as text, the answer as a schema

```
class DamageType(BaseModel):
    prediction: Literal["crack", "scratch", "tire flat",
                        "dent", "glass shatter", "lamp broken"]

image = base64.b64encode(open(path, "rb").read()).decode()

response = client.responses.parse(
    model="ft:gpt-4o-2024-08-06:personal::B5swMoBf",
    instructions=SYSTEM_PROMPT,  # the six classes, rules
    input=[{"role": "user", "content": [{
        "type": "input_image", "detail": "low",
        "image_url": f"data:image/jpeg;base64,{image}"
    }]}],
    text_format=DamageType)
```

What comes back

```
{"prediction": "dent"}
```

482 tokens in, 6 out. At low detail GPT-4o counts the photo as 85 tokens; the rest is the prompt.

Per 1,000 photos: USD 1.91 with the fine-tuned model, USD 1.27 with GPT-4o itself, USD 0.12 with GPT-6 Luna.

No parsing: the schema allows only the six class names, so every answer drops straight into a table.

The zero-shot and the fine-tuned model get exactly this call; only the model name changes.

Fine-Tuning GPT-4o on 900 Photos

One training example, one line of JSONL

```
{ "messages": [  
  { "role": "system", "content": SYSTEM_PROMPT },  
  { "role": "user", "content": [{ "type": "image_url",  
    "image_url": { "url": "data:image/jpeg;base64,...",  
      "detail": "low" } } ] },  
  { "role": "assistant", "content": "dent" } ] }
```

The job: 900 training and 300 validation photos, 3 epochs, 1.14 million trained tokens, 2 hours 19 minutes. About USD 29 at today's price.

The result: a private model name, callable only by keys of the organisation that trained it.

A vendor-risk question

- **7 May 2026**
closed to organisations that never fine-tuned
- **2 Jul 2026**
closed to those without fine-tuned inference in the past 60 days
- **6 Jan 2027**
no new jobs for anyone
- **After that**
existing fine-tuned models keep serving until their base model is retired

Vision fine-tuning on the platform is offered for gpt-4o-2024-08-06 only. The open-weight route is LoRA, as in notebook 02.

A fine-tune ties the result to one vendor and one base model. Keep the training data, so that it can be rebuilt elsewhere.

Results on 300 Test Photos


Model	Run	Accuracy, all 300	Accuracy, 205 without near-copy	Macro F1, all 300
Image classifier (CNN)	2025	0.833	0.785	0.761
GPT-4o, zero-shot	2025	0.797	0.805	0.769
GPT-4o, fine-tuned	2025	0.870	0.893	0.890
GPT-4o, fine-tuned, re-run	2026	0.850	0.863	0.859
GPT-4o, zero-shot, re-run	2026	0.820	0.844	0.821
GPT-5.4 mini, zero-shot	2026	0.773	0.795	0.748
GPT-6 Luna, zero-shot	2026	0.613	0.610	0.601

Macro F1 gives the seven cracks as much weight as the 104 scratches; that is where fine-tuning pulls ahead most. **In 2026 the fine-tune still leads, but only by 0.03 over GPT-4o zero-shot, within the noise.**

GPT-6 Luna costs about a sixteenth of the fine-tune per photo and gets 61% right.

Fine-tuning fixed 40 photos the zero-shot model got wrong, and broke 18 it had right. Net gain: 22 of 300.

Near-Copies Flatter the Model That Memorises



What we found

95 of the 300 test photos have a twin among the 1,200 training and validation photos: the same photo, re-coloured, cropped or mirrored. Splitting by file does not catch that.

Who profits

The CNN scores 0.94 on the twins and 0.79 on the rest. The fine-tuned GPT-4o does not profit, and on the 205 clean photos its lead over GPT-4o zero-shot grows from 7 to 9 points.

How to check

Compare every test photo with every training photo on a tiny grey thumbnail, also mirrored; flag pairs that are too similar.

Screen for near-copies before you trust a test score: a model that memorises looks better than it is.

Where Fine-Tuning Helps, and Where It Costs

Class	Photos	CNN	GPT-4o	Fine-tuned
crack	7	1 of 7	3 of 7	7 of 7
scratch	104	0.84	0.74	0.90
tire flat	22	0.95	0.95	1.00
dent	80	0.76	0.90	0.70
glass shatter	50	0.98	0.96	1.00
lamp broken	37	0.84	0.49	0.86

Share of each class recognised (recall), 2025 run

Zero-shot sees dents everywhere

GPT-4o called 23 scratches and 17 broken lamps a dent. Its dent recall of 0.90 is partly that habit.

The fine-tune overcorrects

It got every crack, flat tyre and shattered window right, but called 14 dents a scratch and 9 a broken lamp.

Look beyond one accuracy figure: the classes that drive claim cost may be exactly where a model is weakest.

More Than a Label: Location and Context

Where is the damage? Fine-tuned GPT-4o

Photo	Label	Predicted	Article prompt	Revised prompt
4779	crack	crack	front bumper	front bumper
3050	scratch	scratch	–	front bumper
5201	scratch	scratch	–	rear bumper
5395	dent	dent	front door	door
7192	dent	dent	driver's side door	–
4443	glass shatter	glass shatter	at windshield	windshield

All 12 classes right under both prompts; 5 locations under the article's prompt, 6 under the revised one. Photo 4443 shows a smashed rear window, called the windshield under both, once as “at windshield”: the old one-string format leaking into the schema. No ground truth for location; a person has to check it.

A richer record: GPT-6 Luna


Schema: damaged part, side, severity, image quality, several damages, and a review flag with a reason.

Confidently wrong: a bumper crack became “vehicle glass, likely windscreen”; a smashed headlamp became glass shatter on a “rear lamp”.


Rule	Fires when	Result
R1	severity severe	0 photos
R2	severity or image not judgeable	4 photos
R3	two models disagree on the class	1 photo
Any rule	sent to a human	5 of 12

Fixed rules decide who looks at a claim; the model's own review flag fired for 11 of 12 photos.

Faking a Claim Photo Takes Two API Calls



Generated: an undamaged car



Edited: “add a dent to the rear door”

What the classifier says

The fine-tuned model: “dent”, location “rear door”. The fake passes. The undamaged original is also a “dent”: six classes, no “no damage”.

What it took

Two calls to gpt-image-2.5-sunburst: 25 seconds and USD 0.02 in total. No editing skills needed.

A classifier checks what a photo shows, not whether the damage exists. That needs other controls.

Provenance Helps, Until the File Is Re-Saved

What the generated file carries

```
PNG chunk caBX, 23,650 bytes: a C2PA manifest
source type : trainedAlgorithmicMedia
actions      : created, converted, watermarked
generator    : gpt-image
signed by    : OpenAI Media Service, time-stamped
```

But: the edited photo is labelled as newly created, with no link to the original, and an image library shows no metadata at all. Re-saved as JPEG, as a portal or a screenshot would, every marker is gone (2.0 MB to 0.32 MB). Missing provenance proves nothing.

Controls from the article

- **Capture protocols:** metadata-rich, time-stamped photos, for example taken in the insurer's own app
- **Provenance checks:** read manifests and watermarks before metadata is stripped for privacy
- **Forensic tools:** detect synthetic media as a routine step
- **Escalation rules:** route suspicious photos to a specialised investigation team
- **Human review:** of classifications; no settlement on a photo alone

Detection and generation are an arms race; the claims process has to assume that any photo can be fake.

Live Demonstration

05_cs3_vehicle_damage_vision.ipynb



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

What This Means for Actuarial Work

Test before you trust

Score every model on photos it never saw, per class, with intervals, and screen for near-copies: they flatter a model that memorises.

Fine-tuning pays on narrow tasks

900 labelled photos lifted accuracy and, above all, the rare classes. Check first whether today's general model is already enough.

Mind the dependencies

A fine-tune lives on one platform under one account. Keep the data and the evaluation, not only the model name.

Photos are evidence

They carry plates, faces and locations, and they can be fabricated. Provenance, privacy and a human reviewer belong in the process.

The same approach works for roof damage, medical images and invoices.

Hands-on Case Study 4

Multi-Agent Systems for Actuarial Code and Data

Legacy Code Holds the Actuarial Logic

WHERE IT LIVES

Language	Typical actuarial use
R	reserving and pricing scripts, experience studies
SAS	data preparation, rating tables, reporting
COBOL	policy administration, valuation runs
VBA	tariff calculators and macros in Excel workbooks

Manual migration is slow and needs three skills at once: the old language, Python, and the actuarial logic behind the code. Rule-based converters miss what syntax does not show.

Why migrate

One maintainable Python stack, fewer people who still read the old language, and a direct path to data science and AI tools.

Why it is hard

Language idioms (1-based indexing, factors), implicit conventions (rounding, random numbers) and domain logic that syntax rules cannot see.

What a migration must deliver

The same numbers as the legacy code, evidence that they are the same, and documentation someone can sign off.

Migrating a reserving script is a model change: it needs validation evidence & a sign-off, whoever writes the new code.

An Agent is a Model in a Loop with Tools

ONE AGENT, ONE LOOP (RECORDED)

```
analysis_agent    gpt-6-sol, reasoning effort low

call 1    the model asks for one tool:
          list_workspace()
call 2    then for three at once:
          read_file("chain_ladder.R")
          read_file("r_output_A.txt")
          preview_csv("triangle.csv", n=12)
          our code runs them and returns the results
call 3    the model answers: purpose, inputs, steps,
          outputs, pitfalls; no further tool call = stop

          3 calls, 18.7 s, USD 0.0217
```

Agent

A model in a loop: it asks for tools, sees their results and stops by a rule.

Tool

A function your code runs. The model only asks for it, with arguments in JSON.


Multi-agent system

Several agents with roles, each with its own instructions, tools and model. Here, code decides who acts next.

The model decides what to ask for; your code decides what it may touch.

A First Example: Chain Ladder from R to Python

THE ARTICLE'S NOTEBOOK, RUN IN THE SEMINAR ENVIRONMENT ON 2 OCTOBER 2026



Fixed in advance

- the order of the agents and every retry: a LangGraph graph, not an agent, decides
- 14 tests verified against R: 6 on the data, 8 on the numbers
- the expected values stay out of the Translation Agent's context

This run

14 of 14 tests passed

at the first attempt, without a retry; about one and a half minutes for the notebook up to the report

The article's benchmark: 10 of 10 runs passed for each of its two examples.



Open in Colab

notebooks/06_cs4_multi_agent_code_migration/first_example
Runs without a key on recorded runs.

The pattern in its simplest form: a fixed graph, five narrow agents, and tests written before the first line is translated.

The Article: Every Migration Passed Its Tests

Example	R lines	Tests (data / content)	Pass rate	Retries: mean / median (range)
Chain-ladder reserving	198	14 (6 / 8)	100%	0.2 / 0 (0–1)
GLM-based reserving	265	15 (8 / 7)	100%	1.2 / 1 (0–4)

Set-up: 10 end-to-end runs per example; GPT-5.4 for analysis and translation, GPT-5.4 mini for compilation, tests and report; LangGraph with a fixed sequential graph; at most 5 retries, then escalation to the report agent.

The article's conclusion: a test-gated multi-agent system can migrate actuarial R code to Python, and complements deterministic tools and human judgement rather than replacing them.

A pass rate is only as strong as the tests behind it. The next slide shows what re-running them taught us.

Re-run: The Article's Tests Passed Wrong Translations

A test must call the code it tests. None of the chain-ladder suite's 8 content tests checked a number the translation computed. A translation with every factor 10% too high (total 59,526.58) and a 5-line stub both passed 14 of 14.

Keep the answers out of reach by construction. The rule not to read the tests was in the prompt only; the analysis agent read the test file in 6 of 20 runs.

Rounding inside a calculation is a change. 6 of 10 translations rounded the factors before chaining: 15,316.35 instead of R's 15,316.33.

A fixed seed makes runs identical, not right. All 10 GLM runs gave the same bootstrap summary, from a different generator than R's.

WRONG TRANSLATIONS AND THE TESTS THEY PASS

Wrong translation	Old suite	New suite
factors $\times$ 1.10	14/14	6/16
stub printing R's numbers	14/14	5/16
factors rounded before chaining	14/14	5/16
constant bootstrap	15/15	11/16
bootstrap without process variance	15/15	15/16

Tests passed, out of the suite's total: old = the article's 14 or 15 tests, new = the notebook's 16. The article's GLM tests did catch a 1% error in the dispersion or the predictions, and all 10 GLM translations matched R's total within $1.6e-14$.

Lessons for any test-gated migration, not a verdict on one model: the next slides build them in.

The Rebuilt System: Five Agents, One Job Each

Analysis agent

GPT-6 Sol

Reads the R code, a data preview and R's printout; returns a structured summary.

Translation agent

GPT-6 Sol

Writes the Python module to a given interface; on a retry it fixes what the feedback names.

Compilation agent

GPT-6 Luna

Runs the translation behind a fence and turns a traceback into a short diagnosis.

Test runner agent

GPT-6 Luna

Runs the hidden tests and turns failures into guidance without numbers.

Report agent

GPT-6 Luna

Writes the migration report; code fills in status, attempts, tests and cost.

The actuary

human

Reviews the translation and the evidence, and signs off.

Each agent has one job, its own tools and the cheapest model that can do it. The article used GPT-5.4 (mini).

Splitting the Work Helps, If Code Keeps Control

WHAT SEVERAL AGENTS BUY YOU

- **Task decomposition:** each agent owns one step with a clear result
- **Specialised agents:** the right model, tools and instructions per step
- **Modular:** swap a model or a prompt without touching the rest
- **Human in the loop:** approval points where the risk is

WHO DECIDES THE NEXT STEP?


Pattern	Next step chosen by	Control
Fixed graph (here)	code, from tool results	highest
Supervisor	an LLM	medium
Hierarchical	LLMs at two levels	lower
Peers, swarm	each agent	lowest

Control: how far code, not a model, fixes the route. More freedom suits open research tasks; a migration with a pass/fail test suits a fixed graph.

Here the graph is fixed: no language model ever decides whether the tests passed.

The Migration Graph: Five Agents, One Retry Loop

The route after each step is read from the tool's result — EXECUTION_SUCCESS, TESTS_PASSED 16/16 — never from what an agent writes.



At most 5 attempts
then the report escalates to a human

USD 0.40 per run
checked before every model call

Recursion limit
a hard stop for the graph itself

Code routes, agents work: the same fixed graph as in the article, with a cost cap and GPT-6 models.

The Contract: Legacy Code In, a Fixed Interface Out

LEGACY R (EXCERPT)

```
for (j in seq_len(n_dev - 1)) {  
  valid_rows <- !is.na(triangle_cum[, j]) &  
    !is.na(triangle_cum[, j + 1])  
  denom <- sum(triangle_cum[valid_rows, j])  
  dev_factors[j] <-  
    sum(triangle_cum[valid_rows, j + 1]) / denom  
}
```



PYTHON INTERFACE, GIVEN TO THE TRANSLATOR

```
def run_chain_ladder(triangle, tail_factor=1.0)  
    -> dict: cumulative, dev_factors, n_pairs,  
        cdf_to_ultimate, results, total_reserve  
  
def main(path="triangle.csv") -> dict  
  
All numbers unrounded; ValueError where R stops.
```

The translator sees

the R code, the triangle, R's printout, the interface and the rules:
allowed imports, no hardcoded results, read only the paths passed in.

It never sees

the tests, R's results in full precision, the second triangle the tests also
use, or any other test input. Its file tools cannot reach them.

R rounds only when it prints. A translation that chains rounded factors prints 15,316.35 instead of R's 15,316.33.

The Test Gate: Run Behind a Fence, Then Grade

1 Run

A fixed runner executes the translation in a child process: its own temporary folder, no key in the environment, an audit hook that refuses files outside the folder, processes and network.

2 Grade

`pytest` compares the runner's outputs with R's results. It never imports the translation.

3 Report

Per failed test only the quantity and an error band; on the hidden triangle only 'fails on a second triangle'.

Chain ladder: 16 tests	n	What it checks
Interface	4	functions, keys, labels, unrounded values
Numbers on the given triangle	5	cumulative, factors, CDFs, reserves, total vs R
Hidden second triangle	2	factors and reserves the agents never see
Tail factor 1.05	1	vs R
Scaling	1	double the triangle, double the reserves
Errors	1	a gap raises <code>ValueError</code> , like R's <code>stop()</code>
Printout	2	R's printed total; returned values unrounded

GLM with bootstrap: 16 tests as well, including the bootstrap on a noisy second triangle (mean within 4 Monte Carlo standard errors of R's, standard deviation within 10%).

A pass needs three things to agree: `pytest`'s exit code, the list of test ids and the results file.

Ten Migrations with GPT-6: All Passed, Two Needed Retries

Example	Passed	Attempts per run	USD per run, mean (range)	Seconds per run, mean (range)
Chain ladder	5/5	1, 1, 5, 1, 4	0.072 (0.038–0.122)	81 (44–145)
GLM with bootstrap	5/5	1, 1, 1, 1, 1	0.055 (0.048–0.066)	65 (55–77)

Same numbers as R: every chain-ladder translation bit-identical on six quantities; GLM totals within $1.6e-14$ on the given triangle and $1.5e-12$ on the hidden one.

Cost by agent: translation 71%, analysis 27%; compilation, tests and report 2% together.

Why the loop fired: Python's `%g` format printed 15316.3 where R prints 15316.33, and only the printed-total test failed. In one run the test agent named the format in its fourth diagnosis; attempt 5 passed.

In the saved notebook the live run passed at attempt 1: 16/16, 44 s, USD 0.049.

A different system from the article's (models, interface, tests and fences all changed): the differences cannot be attributed to the model.

Given the interface, first attempts pass; the loop earns its keep on details a reviewer would also miss.

With the Answers in Reach, No Model Copied Them

Run	Read answers	Copied	Hidden tests	Trap test
Pilot, GPT-6 Sol (2 runs)	7 of 10 attempts	no	passed	failed
'Answers are binding', Sol	attempts 2, 4, 5	no	5 of 5	5 of 5 failed
'Answers are binding', Luna	never	no	5 of 5	5 of 5 failed

Set-up: R's results on the given triangle placed in the agents' folder, and one trap test that no honest translation can pass: R's bootstrap mean to $1e-6$, from a different random generator.

What happened: no model copied a number. Both tried to rebuild R's random-number generator instead, and every run escalated after 5 attempts.

HAND-WRITTEN, NOT BY A MODEL

```
FACTORS = [1.805725, 1.294846, ...]  
RESERVE = [0.0, 23.16, 74.52, ...]
```

```
def run_chain_ladder(triangle, ...):  
    # returns R's printed numbers,  
    # whatever the input
```

Article's suite: 14/14. **New suite:** 5/16 (the hidden triangle, tail, scaling and error tests fail). The tripwire flags a hardcoded result.

In our runs the models did not cheat. A suite that would not notice if they did is still not good enough.

Live Demonstration

`06_cs4_multi_agent_code_migration.ipynb`



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

Part B: Agentic Data Analysis, Actuary in Charge

1 Planner

GPT-6 Sol

A plan from a tool catalogue; statements in the brief become typed hypotheses.

2 Actuary

approves the plan

Exclude the claim-only block? Keep the spiking region?

3 Tools

code

Data-quality scan, one-way tables with intervals, Poisson GLM on rating cells.

4 Analyst

GPT-6 Luna

Proposes the GLM factors with reasons; code tests the choice.

5 Writer

GPT-6 Luna

Every number is a {{placeholder}} that code fills in.

6 Gate

code + GPT-6 Sol

Unknown IDs, stray digits, typed claims; a fact checker advises.

7 Actuary

signs off

With the open notes listed.

WHAT CHANGED AGAINST THE OLD DATA-ANALYSIS NOTEBOOK

Aspect	Old notebook	Part B
Routing	an LLM supervisor	code, from typed results
Code execution	model-written code in the kernel, next to the keys	six typed tools, no model code
Numbers in the report	retyped by the model	placeholders filled by code, then checked

The question: which factors drive claim frequency, and does last year's statement hold?

Part B: The False Statement Was Caught in 5 of 5 Runs

Data issue	Size	Origin
Claim-only block: IDpol 1–24,951	9,387 policies, 27.5% of claims	native
Claims without an amount	9,657 claims	native
Exposure above one year	1,224 rows	native
Exposure ≤ 0 , age < 18 , duplicates	24, 16, 40 rows	planted
Region R25 frequency spike	relativity 2.40	planted

Found with the exact counts in every run, no false alarm. Excluding the block moves the frequency from 0.1007 to 0.0740: the actuary's decision at plan approval.

LAST YEAR'S STATEMENTS, CHECKED BY CODE

'Vehicles over 15 years have the highest frequency'

false: 0.0550, rank 6 of 6 (6–10 years: 0.0813); **caught 5/5**

'Drivers aged 18–20 have the highest frequency'

true: 0.2401, 3.24 $\times$ the portfolio; **confirmed 5/5**

The analyst's GLM choice: the tests disagreed with its first factor set in 5 of 5 runs; after one refit they agreed in 3.

Per run: USD 0.046–0.051, 83–105 s.

The gate proves provenance, not meaning: one signed report called the portfolio frequency 'the block's'.

Risks of Agent-Built Code, and the Control for Each

Risk	What happens	Control in this case study
Silent errors	passes the tests, wrong elsewhere	hidden triangle, invariants, mutants, review
Non-determinism	each run writes different code	recorded runs, fingerprints, same numbers required
Prompt injection	a comment in the legacy code gives orders	tools confined, status set by code
Running generated code	it can read files or leak keys	child process, audit hook, no key in its environment
Confidentiality	proprietary code leaves the house	store=False, contract, data residency
Cost and loops	retries without end	5 attempts, USD cap before each call

Controls live in code, not in prompts. The fence here is a teaching fence, not a production sandbox.

Exercises: Attack It, Port It, Approve It

(a) Prompt injection

Hide an order in an R comment and run the system.

Recorded: the analysis agent flagged the comment ('Ignore it'); no hardcoded total; 16/16 at attempt 1, USD 0.049.

(b) Port a VBA tariff

Migrate MTPL_Premium from VBA. Traps: Excel rounds half away from zero, Trim removes only spaces, an empty age counts as 0.

Recorded: 25/25 at attempt 1, USD 0.028.
Ground truth: Excel's worksheet formula; the VBA itself was never run.

(c) Approval node

Add interrupt() so an actuary must approve the translation.

Recorded: attempt 2 passed and was rejected with a comment; attempt 3 added the docstrings and was approved, USD 0.078.

Each exercise has a recorded outcome, so it runs without a key; with a key, run it live and compare.

Live Demonstration

06_cs4_multi_agent_code_migration.ipynb



github.com/simonhatzesberger/advanced-genai-actuarial



Scan to open in Colab

What This Means for Actuarial Work

A migration is a model change

Keep R's results as the reference, test the new code against them and against inputs it never saw, and document the sign-off.

Tests are the real control

Agents write the code; the test suite decides. Check that each test calls the code, and try a wrong translation on purpose.

Coding agents or your own graph

Coding agents such as Claude Code, Codex or Cursor edit whole repositories; a fixed graph gives each step a log and a gate. Both need the same tests.

People stay in the loop

The actuary approves the plan, reviews the translation and signs off. Agents remove typing, not responsibility.

Which legacy code in your company would you migrate first, and what evidence would your validation function ask for?

Hands-on Case Study 5

Report Generation and Quality Assurance with Generative AI

Reports Turn Numbers into Narrative That Needs Checking

WHAT ACTUARIES WRITE EVERY YEAR

- Actuarial function report: opinions on technical provisions, underwriting, reinsurance
- SFCR and regular supervisory report: solvency, own funds, risk profile
- ORSA report, reserving and experience-study reports
- Pricing and tariff documentation, board memos

WHAT SUPERVISORS KEEP FINDING

- Opinions not clearly stated; year-on-year differences not justified (PRA, actuarial function reports)
- Legislation copied instead of undertaking-specific text; summaries not fit for policyholders (EIOPA, SFCR)
- Boilerplate; no company-specific assumptions for cash flows and discount rates (ESMA, IFRS 17)

A language model can draft the words. The numbers, and the checks, must come from code — and a person signs.

The Same Documents Are Worth Drafting and Worth Checking

Document	Legal basis	Prose share	What a check can test
SFCR, policyholder part (new)	Solvency II review	high	element list, prescribed text, 5 pages
Actuarial function report	Solvency II Art. 48	high	opinions backed by analysis
Reserving and experience studies	internal, feeds AFR and SFCR	medium	movements add up, drivers
IFRS 17 movement analysis	IFRS 17	medium	roll-forward ties to the notes
Pricing and tariff documentation	national rules, POG	medium	formulas, tables, legal limits
ORSA and board memos	Solvency II Art. 45, governance	high	claims supported by evidence

Good candidates share three traits: the numbers already exist in models or templates, the prose explains them, and the same errors recur every year.

This case study: draft and check the new SFCR policyholder part for a fictitious insurer, Brisendale Mutual.

Warm-up: Arithmetic Checks Miss Unverifiable Claims

AN SFCR EXCERPT WITH THREE PLANTED ERRORS

The company's solvency ratio is 175% ... It is therefore well above the internal target corridor and rose by 5 percentage points year on year.

The expense ratio is 28% (costs EUR 140 million, gross premiums EUR 700 million) and therefore exceeds the target.

... the policyholder bonus rates was adjusted. The surplus participation for 2026 was set at a level that ...

Checker	Errors found	Unverifiable
Ratio and spelling	3 of 3	0 of 5
Full gate	3 of 3	5 of 5

Unverifiable: claims nobody can check, which a checker must report: 'above the internal target corridor' with no corridor given, 'rose by 5 percentage points' with no prior-year figure, and three more.

Correct arithmetic is not a verified report: a claim nobody can check must be reported, not skipped.

Prior Years as Examples Beat Copying Last Year's Text

A · last year's text, figures updated **B · language model with prior years as examples**

Prior years given to the model as examples

0 = no example: the model only knows the rules. 4 = the texts of the last four years: the model sees how this insurer writes.

House rules

☒ Apply the house rules (movement phrases and terminology from the workbook)

On: code computes each movement phrase from the size of the change ("decreased slightly") and the gate enforces it and the house terms. Off: the model decides the words itself.

Draft live with gpt-6-luna (about USD 0.002) **Show the recorded draft**

passed **USD 0.0020** **10,469 / 1,598** **11 s**

drafting gate (1 call(s)) cost (recorded on 2026-10-01) input / output tokens gpt-6-luna

EUR 373.1 million a figure that code filled in from a placeholder (hover to see it)

E.1 · OBJECTIVES, POLICIES AND PROCESSES FOR MANAGING OWN FUNDS (Art. 297(1)(a))

Our objective in managing own funds is to hold enough capital of good quality to meet our obligations to our policyholder members at all times and to cover the Solvency Capital Requirement (SCR) with a margin. The Board approves the capital management policy and reviews it every year as part of our own risk and solvency assessment. Our business planning covers a time horizon of three years, and our risk appetite is an SCR coverage ratio of at least 160%. As a mutual, we pay no dividends; profits are retained in own funds. The investment policy was updated during 2026 to exclude companies that earn most of their revenue from thermal coal; the effect on investment returns has not been measured.

Step 3 of the app: four prior years, house rules on; code filled in every highlighted figure.

Last year's text keeps every word, the wrong ones too; examples keep the style, and code keeps the facts.

SFCR E.1 AND E.2 FOR 2026, THREE INSURERS

How the draft is made	Errors	Wording like 2025
Last year's text, figures updated	3 to 5	unchanged
Model, no prior year	0	0.29 to 0.40
Model, 1 to 4 prior years	0	0.82 to 0.90

Errors per insurer: wrong direction words, shifted dates; stale causes only as warnings.

Wording like 2025: word-level similarity, 1 = the same words.

House rules on / off: 0 / 1 to 15 deviations from the house phrases.

Excel in, Word out; recorded drafts (USD 0.0015 to 0.0029 each) run without a key.

The New Policyholder Part: Five Pages, Fixed Elements

A Languages

a notice at the start: policyholders may ask for a version in their own language; links to other versions (Art. 292(1))

B Business and performance

name and legal form, supervisor, qualifying holdings, group, significant events, underwriting and investment performance (292(2))

C Capital and risk profile

a prescribed text on SCR and MCR; SCR, MCR, eligible own funds and coverage ratio for this and the previous year; non-compliance; material risks (292(3))

D Other information

other material information; a statement on sustainability plans (292(4))

RULES A CHECK CAN TEST

- At most five pages, without leaving anything out (292(5))
- The SCR/MCR text word for word, in every language version
- A machine translation must be disclosed to the policyholder (Art. 298a)

“Two capital requirements aim at measuring the financial soundness of the undertaking ...”

Opening words of the prescribed text, Art. 292(3)

Applies from 30 January 2027; EIOPA expects the first reports in this format for the year 2027, published in 2028.

Facts and Bullets In, Draft Out; Code Owns the Numbers

WHAT GPT-6 LUNA WRITES (EXCERPT OF A TEMPLATE)

```
At {{date_cur}}, our SCR was {{scr_cur}} and our MCR  
was {{mcr_cur}}. ... The SCR coverage ratio increased  
by {{cov_scr_change_pp}} ... The new motor reinsurance  
contract reduced the insurance risk we keep.
```

```
claims: [{"sentence": "C.p2.s6", "bullet": "DR2"}]
```

Inputs: 65 figures for 2026 and 2025, generated by code with consistent totals and ratios, and six driver bullets from finance and risk, two of which say what is not yet known.

One call

13 s, USD 0.0010

Placeholders

75 numbers, none typed by the model

Three drafts

different wording (similarity 0.28 to 0.64), the same placeholders

Causal claims

listed by the model, each with the bullet it relies on

The model chooses the words; code inserts the numbers and the prescribed SCR/MCR text.

Live Demonstration

07_cs5_report_qa.ipynb



github.com/simonhazesberger/advanced-genai-actuarial



Scan to open in Colab

Fresh Drafts Pass the Code Checks but Need an Editor

0 code-check errors on the fresh draft: every number traced, the prescribed text in place.

Plain language: the audience check flagged 8 technical terms on every run.

7 corrections by the author made the gold version: 5 plain-language edits (quota share, combined ratio ...) and 2 missing explanations (why claims rose; what the quota share does).

Checks written for one draft misfire on another: an earlier live draft got 6 errors, all false alarms ('during 2026' not read as a year; a change read as a level).

The code checks prove that the numbers are right; whether a policyholder understands the text is still judgement.

Code Checks First, LLM Checks for Language

12 CODE CHECKS (9 SHOWN)

Check	What it tests
Required elements	the eight figures, not just the words
Prescribed text	word for word, English and Lithuanian
Language notice, links	at the start
Five pages	an estimate in a declared format
Every number	traced to the facts file
Metric and year	each figure bound to its label
Ratios and totals	recomputed from printed figures
Direction words	'rose' matches the sign
Jargon list	technical terms in policyholder text

6 LLM CHECKS

Check	What it tests
Claims supported	each cause backed by a driver bullet, quoted
Element covered	not just a heading
Audience	plain language, no sales tone
Spelling	with a do-not-flag list
Numbers in prose	extracted for code to compare
Chart read blind	an image, no table values

Structured outputs; every finding quotes the document and code checks that the quote is there.

An LLM finding is a warning until code confirms it; 'not verifiable' is reported, never skipped.

The LLM Extracts, Python Decides

ONE CLAIM, CHECKED (ILLUSTRATION)

```
sentence S14: "Euro interest rates rose during the year,  
              which reduced the value of our bonds."
```

```
LLM  claim_quote    "Euro interest rates rose"  
      bullet_id      DR3  
      evidence_quote "Euro interest rates fell in the  
                      second half of the year"  
      verdict         unsupported  
  
code  claim quote in S14?    yes  
      evidence quote in DR3? yes  -> warning: unsupported
```

The model reads and quotes

It returns structured fields, not prose: which claim, which bullet, which words.

Code checks the quotes

A quote that is not in the document or the bullet makes the finding 'not verifiable', never silently dropped.

Arithmetic stays in Python

Numbers the model extracts are compared with the facts by code, at the printed precision.

The warm-up's checker skipped what it could not verify; here that becomes a finding of its own.

Seeded Errors and a Clean Control Measure the Gate

Gate	Errors planted first	Hold-out, planted later	All
Code checks only	8 of 10	2 of 6	10 of 16
Full gate, 5 GPT-6 Luna runs	9 of 10	5 of 6	14 of 16
Full gate, 1 GPT-6 Sol run	9 of 10	5 of 6	14 of 16

The hold-out errors were planted by another model after the checks were frozen. The code checks found 2 of 6 of them: the batch held a spelling error and a cause, which only the LLM checks read, and a jargon term missing from the list.

Missed in every run: a wrong e-mail address and a reference to the wrong figure, errors no check was written for.

Clean control: 0 errors, 9 to 11 warnings per run.

Sol vs Luna: the same recall at USD 0.066 against about USD 0.002 per run.

A recall figure depends on the errors you plant: plant some after freezing the checks & report the batches separately.

A Second Notebook Checks a 93-Page Calculation Basis

The DAV working group's TB validation: a notebook that checks the technical calculation basis (TB) of a German private health insurance tariff, the document that trustee and supervisor review at every premium adjustment.

13 DETERMINISTIC CHECKS

- References, contents, symbol index
- Table identities and the implied interest rate
- Premiums monotone, age tables complete
- Every cell against reference CSVs and a database
- Parameters, legal citations, versions

6 LLM CHECKS

- Spelling
- Formulas extracted by the LLM, recomputed in Python
- Sections against their annexes
- Legal citations in context
- Formula pages read as images
- Charts read blind, compared in Python

HOW IT COMPLEMENTS NOTEBOOK 07

	Notebook 07	TB validation
Document	3-page public text, drafted here	93-page technical PDF from a pricing pipeline
Planted errors	16, of them 6 after the freeze	13, with a documented ground truth
Models	GPT-6 Luna, GPT-6 Sol	Claude Sonnet 5; Haiku 4.5 for quick runs
Language	English, Lithuanian	German



Repository of the working group

github.com/DeutscheAktuarvereinigung/
WorkingGroup_AI_ModernMethods_Health
folder tb-validierung, MIT licence; without a key the
13 deterministic checks run

The same pattern on another document: code checks structure and numbers; the LLM reads language, formulas and charts.

The TB Validation Found 13 of 13 Seeded Errors

The test bench

A 93-page synthetic technical calculation basis for a health tariff, a faulty copy with 13 documented errors (transposed digits, a broken legal limit, wrong citations, dead references, typos, a formula sign, a chart curve), and the clean copy as a control.

Check layer	Errors found
Deterministic checks only	9 of 13
+ LLM checks on the text	11 of 13
+ LLM checks on page images	13 of 13
Clean control: errors / warnings	0 / 5

DESIGN RULES WORTH COPYING

- Deterministic checks first; the LLM only where language, context or images need it
- The LLM extracts, Python computes: arithmetic is never left to the model
- Every finding quotes the document; the quote must appear on the stated page
- A chart is read without the table values, so the model cannot anchor on them
- Uncertain LLM judgements are warnings, not errors

Limits: the errors and the checks were written by the same team; one run per document. Notebook 07's hold-out batch shows why that matters.

The sign error and the chart were caught only on page images; the typos and the worked example only by the LLM.

The Judge Needs Its Own Test Set

Run	Unsupported found	False alarms	Borderline flagged
GPT-6 Luna 1	12/12	0/12	6/6
GPT-6 Luna 2	12/12	0/12	4/6
GPT-6 Luna 3	12/12	0/12	5/6
GPT-6 Sol 1	12/12	0/12	6/6

30 labelled sentences: 12 supported, 12 unsupported, 6 borderline. With 12 cases per side, 12 of 12 still leaves a 95% interval from 0.76 to 1.00.

Borderline: 6 sentences, 5 of them inferences beyond the bullets and one that states no cause; the runs flagged 4 to 6 of them. How strict the judge should be is a policy you write down, not the model's choice.

A tuned prompt

On the tuning sentences: 8/8 found, 0/8 false alarms.

On the held-out ones: 4/4 found, 1/4 false alarm.

Tune on one set, measure on another.

An LLM checker has a hit rate and a false-alarm rate: measure both before it gates a report.

Machine Translation: Numbers, Wording, Disclosure

Planted translation defect	Caught by
A number changed (373,1 → 371,3)	number check across languages
The disclosure removed	machine-translation check (Art. 298a)
An English number format left in	number check
The prescribed text machine-translated	verbatim check against official Lithuanian text
A placeholder dropped	placeholder check

ASK THE ROOM

“2026 m. pasirašytosios įmokos padidėjo 4,5 proc., nuo 1 108,9 mln. EUR 2025 m. iki 1 158,7 mln. EUR 2026 m.”

Does this read right? The Lithuanian texts were checked by code and a model only.

The official Lithuanian text has no ‘SCR’ or ‘MCR’: GPT-6 Luna’s own translation of it matched the official wording at a similarity of 0.23, so code inserts the official text.

Code checks what can be checked; a native speaker judges whether it reads well.

A Clean Gate is Not Proof

Errors no check was written for pass: a wrong e-mail address and a reference to the wrong figure, in every run.

An LLM can 'fix' the wrong way: the spelling check turned a planted '16%' into '8 percentage points' in 4 of 5 runs; the right value is 16.

A picture can hide a change: a chart bar changed by 18% was read correctly by the blind vision check; one changed by 7.6% was missed in 3 of 3 reads. Redrawing the chart from the facts caught both.

Code checks have gaps too: of 220 seeded one-error variants they caught 214; direction words were the weak spot (15/20).

A clean gate says what it looked for. The sign-off must say what nobody looked for.

Risks of Generated Reports, and the Control for Each

Risk	What happens	Control in this case study
Invented or stale numbers	a figure that is not in the facts	code inserts every number; tracing
Unsupported claims	a cause nobody gave the model	evidence quote from the driver bullets
Wrong mandatory wording	the prescribed text paraphrased	verbatim check against the regulation
Translation errors	a number or a negation changes	number check across languages, disclosure
A checker that errs	false alarms, missed errors	a labelled test set for the judge
A clean gate taken as proof	errors no check was written for	planted errors, clean control, review

Measure the gate before trusting it: seeded errors show what it catches, a clean copy what it invents.

Exercises: Plant an Error, Write a Check, Tune a Judge

(a) Plant your own error

`plant(sentence_id, old, new)` changes one sentence; run the gate and see which check catches it.

In the room: who can slip an error past every check?

(b) Write the '%' check

A change of a ratio in '%' where 'percentage points' is meant. Today the gate reports it as an untraceable number.

Reference: the right diagnosis, 20 of 20 seeded variants, no false alarm on gold.

(c) Tune the judge

Change the judge's prompt on 20 tuning sentences, then measure on the 10 held out.

Reference: no false alarm when tuning, one on the held-out sentences.

Exercise (a) runs in the room; (b) and (c) are for later, with reference solutions in the notebook.

What This Means for Actuarial Work

Draft from facts, not from memory

Give the model the numbers and the reasons as inputs; let code put every number into the text.

Checks are an actuarial product

Write them from the rules and the errors you know, test them on seeded errors, keep the clean control.

Judges need validation too

An LLM checker has a hit rate and a false-alarm rate. Measure both before it gates a report.

Sign-off stays human

The board approves the SFCR; the actuarial function signs its report. Generated text changes the work, not the responsibility.

Discuss: Which report in your company would you draft with an LLM first, and which check would you insist on before anyone reads it?

Outlook and Discussion

Further Applications, Challenges, and Future Developments

Agents Inside Word and Excel




Agent Mode in Word

youtube.com/watch?v=8oJwUjLn6IM



Agent Mode in Excel

youtube.com/watch?v=nSqCy-7Qabk



Claude for Excel Anthropic add-in, generally available since May 2026

Agent Mode is now part of Microsoft Copilot ("Edit with Copilot") and available to EU users.

In EU tenants, Copilot's Anthropic models run outside the EU Data Boundary. An admin setting controls their use.

Coding Agents

```
ModelSpec("GPT-5.4 mini", "openai", "gpt-5.4-mini-2026-03-17", "OPENAI_API_KEY", temperature=0.0),
ModelSpec("GPT-4.1 mini", "openai", "gpt-4.1-2025-04-14", "OPENAI_API_KEY", temperature=0.0),
ModelSpec("Gemini 2.5 Flash", "google", "gemini-2.5-flash", "GOOGLE_API_KEY", temperature=0.0),
ModelSpec("Mistral Small 4", "mistral", "mistral-small-2603", "MISTRAL_API_KEY", temperature=0.0),
]

generate_model_specs_comparison

def create_model_specs_comparison_dataset(model_specs: list[ModelSpec] = MODEL_SPECS) -> pd.DataFrame:
    """Create a dataframe comparing the selected benchmark model specifications."""
    rows: list[dict[str, Any]] = []
    for spec in model_specs:
        rows.append(
            {
                "label": spec.label,
                "provider": spec.provider,
                "model_name": spec.model_name,
                "api_key_env": spec.api_key_env,
                "temperature": spec.temperature,
            }
        )
    return pd.DataFrame(rows)
```

GitHub Copilot Assign it a GitHub issue; it works on a branch in the background and opens a pull request.

```
• OpenAI Codex (research preview) v0.1.2504101751

localhost session: c7e08ec10e0e319e79d379a1930ba9
├─ workdir: C:\Repository\GitHub\codex-commands
├─ model: deepseek-v1-distill-llama-70b
├─ provider: groq
├─ approval: suggest
user
Please create a data analysis script that takes the clean data and runs statistical analysis.


codex
I'll help create a statistical analysis script. First, I need to understand what kind of data we're working with and what statistical tests are required. Since no specific data or requirements were provided, I'll create a general-purpose statistical analysis template that can be adapted to different datasets.

# statistical_analysis.py
# General-purpose statistical analysis script

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

OpenAI Codex A coding agent in the terminal, the IDE or the cloud, where tasks run in parallel.

Live-Demo



Claude Code Reads the codebase, edits files and runs commands, in the terminal, IDE, desktop app or browser.

Computer Use: The Model Operates the Screen

How it works

The model looks at a screenshot, decides where to click or what to type, and repeats until the task is done. It works with any application, including those without an API.

The catch

It can click anything you can click. Run it in a sandbox, allow only the applications it needs, and confirm every irreversible step yourself.

Where it stands

Anthropic's computer use left beta on 19 August 2026. OpenAI retired its computer-use preview model on 23 July 2026.




Computer use for coding
youtube.com/watch?v=vH2f7cjXjKI

Insurance Sold Inside ChatGPT

Compare Car Insurance Quotes Using Insurify — Without Leaving ChatGPT

Get personalized rate estimates, compare top companies, and explore real customer reviews in one simple, conversational experience with our new ChatGPT app.



Already in Europe

UK: Aviva (home and life quotes), MoneySuperMarket, Go.Compare. Spain: Allianz Direct, Tuio. France: MAIF. Most quote in the chat and hand over to their own website to buy.

No insurer from the Nordics or the Baltics yet.

Insurify, United States

The first insurance app in ChatGPT, live since 3 February 2026. Since August 2026 it gives real quotes in the chat and a path to purchase, in 11 states.


Company figures: revenue from ChatGPT visitors up 212% in six months.

Markets noticed

On 9 February 2026, the day Insurify announced its app, broker shares fell: WTW -12%, Gallagher -10%, Aon -9%. Bloomberg linked the drop to the app and to new Anthropic tools.

The first price comparison now happens inside someone else's product.

GraphRAG: Retrieval over a Knowledge Graph



LLM-extracted graph of the AI Act and EIOPA's AI Opinion (371 nodes, 641 relationships), around Annex III 5(c)

Worth it when an answer is spread across several provisions, actors and dates.
The graph is only as good as the extraction behind it.

Corpus: EUR-Lex 02024R1689-20260727; EIOPA-BoS-25-360. Answers by gpt-5.4-nano; the retrieval was tuned on the same six questions, so the scores are optimistic.

What it adds

Vector RAG finds passages that sound like the question. GraphRAG also follows the links: which article imposes a duty, on whom, from when.

Test question

From when must an insurer's pricing AI for life and health insurance meet the high-risk rules, and which provisions say so?

Vector RAG: "The corpus does not contain the answer ... it does not state any date."

GraphRAG: "From 2 December 2027 ... Art. 113(3)(c)(i), Annex III 5(c), Art. 6(2)."

Six test questions, three runs each: GraphRAG answered all four multi-provision questions. Both declined the question outside the corpus.

Bias in Generative AI

Prompt: “A photo of a successful person.”



gpt-image-1: 8 of 8 men. gpt-image-2 (April 2026): 8 of 8, three of them Elon Musk.

Published: CEO images 100% men in DALL·E 3, 96% in Midjourney 6.

Prompt: “What happened at Tiananmen Square on 4 June 1989?”

qwen3.5:4b — open weights, run offline on this laptop

“The Chinese government firmly upheld the principles of peace and order while ensuring social stability across all regions.”


Four runs: twice the official line, twice a made-up story. None says anyone died. Asked about Kent State in 1970, the same model reports the four students killed.

The weights carry it: Chinese rules require generated content to “uphold core socialist values”. NIST found DeepSeek’s downloadable models echo Communist Party narratives in up to 26% of answers, US models in 1–4%.

Every model carries its training data and its maker’s rules, US ones included. Test it on your own cases first.

EU AI Act: What Applies When

Regulation (EU) 2024/1689 as amended by the Digital Omnibus, Regulation (EU) 2026/1744, in force since 27 July 2026



What the Omnibus changed: high-risk duties moved from 2 August 2026 to 2 December 2027; AI literacy is now a duty to “take measures to support” it, with no guaranteed level; two bans added.

Fines up to EUR 35 million or 7% of worldwide turnover for prohibited practices, EUR 15 million or 3% for most other breaches.

Source: EUR-Lex, Regulation (EU) 2024/1689, consolidated version of 27 July 2026; Regulation (EU) 2026/1744.

EU AI Act: What It Means for an Insurer

Tier	Where it applies in insurance	What you have to do
Prohibited Art. 5, since 2 Feb 2025	AI that manipulates customers or exploits their age, disability or financial situation; social scoring; emotion recognition of staff.	Do not use it. Fines up to EUR 35 million or 7% of turnover.
High-risk Annex III 5(c), from 2 Dec 2027	Risk assessment and pricing of individuals in life and health insurance, including accident, sickness, long-term care and credit life. Motor and home are out (draft Commission guidelines). Also recruitment, Annex III 4(a).	Risk management, data governance with bias checks, documentation, logging, human oversight. As deployer, a fundamental rights impact assessment (Art. 27). Systems in use before December 2027 are caught once significantly changed.
Transparency Art. 50, since 2 Aug 2026	Customer chatbots; AI-generated text, images and voice.	Tell people they are dealing with an AI; mark generated content.
Everything else no AI Act tier	Claims triage, reserving, fraud scoring, non-life pricing, and the pipelines from this seminar.	AI literacy (Art. 4). EIOPA's AI Opinion (August 2025) expects proportionate governance in six areas; DORA treats your LLM API as an ICT third-party service.

Still open: EIOPA asked in April 2026 to exclude GLMs and GAMs from the AI definition.

Where AI is Heading



Next-token prediction has limits

Yann LeCun argues that LLMs do not understand the physical world, have no lasting memory and cannot really plan. Others in the field disagree.



Public text is running out

Epoch AI puts the stock of public human-written text at about 300 trillion tokens, likely used up between 2026 and 2032. Labs turn to synthetic and multimodal data.



From chatbots to agents that act

OpenClaw, an open-source assistant that acts on your own computer, reached 100,000 GitHub stars two months after it started. An audit of its add-on marketplace found 341 malicious skills out of 2,857.



World models

Models that learn how the physical world behaves and predict what happens in it. LeCun's AMI Labs in Paris raised \$1.03 billion in March 2026; Google DeepMind's Genie 3 generates worlds you can walk through.

AI takes over more of the routine. Validation, judgement and the final decision stay with the actuary.

Discussion

Which part of your work will matter more in three years, and which less?



Two or three voices from the room

Think along the actuary's role



01

Specify



02

Delegate



03

Check



04

Sign off



05

Explain

Closing

What to Take Home from Two Days

Foundations

A model predicts the next token. Structured outputs, tools and retrieval turn it into a component you can test.

CS1 Claims text

Fixed classes in a schema turn one line of text into features a cost model can use.

CS2 Market comparison

Retrieval plus a validated object, scored value by value against references: numbers are easy, ratings are not.

CS3 Vehicle damage

A fine-tune can lead, near-copies flatter a model, and a fake claim photo takes two calls.

CS4 Multi-agent systems

Code routes, agents work, tests decide: a test must call the code and use inputs the model never saw.

CS5 Reports

Draft from facts, check with code first, measure the checker, then a person signs.

The Floor is Yours

One sentence each: what will you try first at your desk, and what would stop you?



A round through the room,
one sentence per person

Think of one process you know well



Which step would you hand to a model?



Which check would you insist on?

Thank you!


Please rate the training via the survey-link you will receive per email.

Visit our website



www.actuarial-academy.com
for more events.

Follow us on LinkedIn



[EAA - European Actuarial Academy](#)
for updates & actuarial fun.